

Q. Does user's demographic(location, age, gender) impact type of anime(genre, episodes, anime type(OVA, movie), Rank) they are prefer? For solving this question, I am trying to determine the rating of an anime based on user's demographic(Age, Location, Gender) and anime's features(genre, type and episodes)

The first algorithm I am trying is OrderedModel. I am trying to classify a given combination of user demographic and anime features to the probable rating by user. This is because rating is a multiclass column. Plain logistic regression would be a bit lacking to predict the expected rating. OrderedModel is more suitable for ordinal multiclass classification.

```
import pandas as pd
import numpy as np
from statsmodels.miscmodels.ordinal_model import OrderedModel
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MultiLabelBinarizer

data = pd.read_csv('../joined_datasets/joined_rating_dataset.csv')
cleaned_dataset =
pd.read_csv("../cleaned_datasets/users_details_dataset_cleaned.csv"
)
data.dropna()
data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5207686 entries, 0 to 5207685
Data columns (total 21 columns):
#   Column                                Dtype
---  -
0   Unnamed: 0                            int64
1   user_id                               int64
2   anime_id                              int64
3   rating                                int64
4   Gender                                object
5   Location                              object
6   Birthday_Date                         object
7   Joined_Date                           object
8   Age_Join                              float64
9   Episodes Watched                      float64
10  Age                                    float64
11  Name                                   object
12  Genres                                object
13  Type                                  object
14  Start Date                            object
15  End Date                              object
16  Studios                               object
17  Source                                object
18  Rank                                  object
19  Episodes                              float64
20  Episodes_Norm                         float64
```

```
dtypes: float64(5), int64(4), object(12)
memory usage: 834.4+ MB
```

Performing undersampling as number of rows was a bit too high to compute using my laptop.  
Considering only the top 10 countries

```
df = cleaned_dataset.copy()
import pandas as pd
import matplotlib.pyplot as plt

location_user_counts = df['Location'].value_counts()

top_countries = df['Location'].value_counts().head(10)

# # Get the list of top 20 countries
top_10_countries = top_countries.index.tolist()
data = data[data['Location'].isin(top_10_countries)]
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 3357313 entries, 0 to 5207685
Data columns (total 21 columns):
#   Column                Dtype
---  -
0   Unnamed: 0             int64
1   user_id                int64
2   anime_id               int64
3   rating                 int64
4   Gender                 object
5   Location               object
6   Birthday_Date          object
7   Joined_Date            object
8   Age_Join               float64
9   Episodes Watched       float64
10  Age                    float64
11  Name                   object
12  Genres                  object
13  Type                   object
14  Start Date             object
15  End Date               object
16  Studios                 object
17  Source                 object
18  Rank                   object
19  Episodes                float64
20  Episodes_Norm           float64
dtypes: float64(5), int64(4), object(12)
memory usage: 563.5+ MB
```

```
# Undersampling
# Count entries for each location
```

```

location_counts = data['Location'].value_counts()

# Decide the target sample size (e.g., use the minimum count of the
top 5 locations)
target_sample_size = location_counts.nsmallest(5).min() # Choose the
minimum of the top 5

# Under-sample the DataFrame
under_sampled_df = data.groupby('Location').apply(lambda x:
x.sample(n=min(len(x), target_sample_size),
random_state=42)).reset_index(drop=True)

# Count occurrences of each class in the target variable (assuming
your target variable is 'rating')
target_counts = under_sampled_df['Location'].value_counts()
# Output the results
print("Under-sampled DataFrame shape:", under_sampled_df.shape)
print("Counts of each class in the target variable:")
print(target_counts)

/var/folders/dz/fg9tl53x4y16ytmghwdt0kr0000gn/T/
ipykernel_85581/807053095.py:9: DeprecationWarning:
DataFrameGroupBy.apply operated on the grouping columns. This behavior
is deprecated, and in a future version of pandas the grouping columns
will be excluded from the operation. Either pass
`include_groups=False` to exclude the groupings or explicitly select
the grouping columns after groupby to silence this warning.
  under_sampled_df = data.groupby('Location').apply(lambda x:
x.sample(n=min(len(x), target_sample_size),
random_state=42)).reset_index(drop=True)

Under-sampled DataFrame shape: (1075590, 21)
Counts of each class in the target variable:
Location
Australia      107559
Brazil          107559
Canada          107559
France          107559
Germany         107559
Philippines     107559
Poland          107559
Russia          107559
Sweden          107559
United States   107559
Name: count, dtype: int64

```

Checking if more undersampling is needed based on ratings.

```

from collections import Counter

```

```

counter = Counter(under_sampled_df["rating"])


for k, v in counter.items():
    per = 100*v/len(under_sampled_df["rating"])
    print(f"Class = {k}, n={v} ({per:.2f}%)")

target_count = under_sampled_df["rating"].value_counts()
target_count.plot(kind='bar', title='Count (target)')

Class = 10, n=142811 (13.28%)
Class = 9, n=193198 (17.96%)
Class = 6, n=123165 (11.45%)
Class = 8, n=264913 (24.63%)
Class = 7, n=243352 (22.62%)
Class = 5, n=61613 (5.73%)
Class = 4, n=25777 (2.40%)
Class = 2, n=5855 (0.54%)
Class = 3, n=10901 (1.01%)
Class = 1, n=4005 (0.37%)

<Axes: title={'center': 'Count (target)'}, xlabel='rating'>

```



```
data.columns
```



4	Birthday_Date	808423	non-null	object
5	Joined_Date	808423	non-null	object
6	Age_Join	808423	non-null	float64
7	Episodes_Watched	808423	non-null	float64
8	Age	808423	non-null	float64
9	Name	808423	non-null	object
10	Start Date	808423	non-null	object
11	End Date	808423	non-null	object
12	Studios	808423	non-null	object
13	Source	808423	non-null	object
14	Rank	808423	non-null	object
15	Episodes	808423	non-null	float64
16	Episodes_Norm	808423	non-null	float64
17	Gender_Male	808423	non-null	float64
18	Gender_Non-Binary	808423	non-null	float64
19	Location_Brazil	808423	non-null	float64
20	Location_Canada	808423	non-null	float64
21	Location_France	808423	non-null	float64
22	Location_Germany	808423	non-null	float64
23	Location_Philippines	808423	non-null	float64
24	Location_Poland	808423	non-null	float64
25	Location_Russia	808423	non-null	float64
26	Location_Sweden	808423	non-null	float64
27	Location_United States	808423	non-null	float64
28	Type_Music	808423	non-null	float64
29	Type_ONA	808423	non-null	float64
30	Type_OVA	808423	non-null	float64
31	Type_Special	808423	non-null	float64
32	Type_TV	808423	non-null	float64
33	Genre_Action	808423	non-null	float64
34	Genre_Adventure	808423	non-null	float64
35	Genre_Avant Garde	808423	non-null	float64
36	Genre_Award Winning	808423	non-null	float64
37	Genre_Boys Love	808423	non-null	float64
38	Genre_Comedy	808423	non-null	float64
39	Genre_Drama	808423	non-null	float64
40	Genre_Ecchi	808423	non-null	float64
41	Genre_Fantasy	808423	non-null	float64
42	Genre_Girls Love	808423	non-null	float64
43	Genre_Gourmet	808423	non-null	float64
44	Genre_Horror	808423	non-null	float64
45	Genre_Mystery	808423	non-null	float64
46	Genre_Romance	808423	non-null	float64
47	Genre_Sci-Fi	808423	non-null	float64
48	Genre_Slice of Life	808423	non-null	float64
49	Genre_Sports	808423	non-null	float64
50	Genre_Supernatural	808423	non-null	float64
51	Genre_Suspense	808423	non-null	float64

```

dtypes: float64(44), object(8)
memory usage: 326.9+ MB

df = df.astype({col: 'float' for col in
df.select_dtypes(include=['bool']).columns})
df = df.astype({col: 'float' for col in
df.select_dtypes(include=['int64', 'float64']).columns})
df = df.dropna()
df = df[df["Rank"]!="UNKNOWN"]

```

Dropping columns with no data

```

genre_columns = [col for col in df.columns if
col.startswith("Genre_")]

genre_counts = df[genre_columns].sum()

zero_count_genres = genre_counts[genre_counts == 0].index
zero_count_genres
df = df.drop(columns=zero_count_genres)

```

Dropping type columns with no data

```

type_columns = [col for col in df.columns if col.startswith("Type_")]

type_counts = df[type_columns].sum()

zero_count_types = type_counts[type_counts == 0].index
df = df.drop(columns=zero_count_types)

```

Standardizing the columns with high variance data, to a small range(0-1)

```

# Standardize the features
features = df.columns.difference(['user_id', 'anime_id', 'rating',
'Birthday_Date', 'Joined_Date', 'Age_Join', 'Episodes Watched',
'Start Date', 'End Date', 'Name',
'Rank', 'Studios', 'Source', 'Episodes_Norm', 'Unnamed: 0'])
from sklearn.preprocessing import LabelEncoder, StandardScaler
scaler = StandardScaler()
df_scaled = scaler.fit_transform(df[features])

# # Split the data
X = df_scaled

```

```

y = df[target]
y = y.astype('category')

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

```

```

# Fit an Ordinal Logistic Regression model
model = OrderedModel(y_train, X_train, distr='logit')
model_fit = model.fit(method='bfgs')
print(model_fit.summary())

```

```

# Coefficient significance
print(model_fit.pvalues)
print(model_fit.conf_int())

```

```

/Users/ramachandrak/Repos/MS/sem1/.venv/lib/python3.12/site-
packages/statsmodels/miscmodels/ordinal_model.py:205: Warning: the
endog has ordered == False, risk of capturing a wrong order for the
categories. ordered == True preferred.

```

```

warnings.warn("the endog has ordered == False, "

```

```

Optimization terminated successfully.
Current function value: 1.809041
Iterations: 76
Function evaluations: 77
Gradient evaluations: 77

```

#### OrderedModel Results

```

=====
=====
Dep. Variable:                rating    Log-Likelihood:    -
1.1700e+06
Model:                OrderedModel    AIC:
2.340e+06
Method:                Maximum Likelihood    BIC:
2.341e+06
Date:                Tue, 05 Nov 2024

Time:                01:05:53

No. Observations:                646738

Df Residuals:                646693

Df Model:                36

```



	coef	std err	z	P> z	[0.025
0.975]					
x1	-0.0982	0.002	-42.858	0.000	-0.103
-0.094					
x2	0.0701	0.003	27.653	0.000	0.065
0.075					
x3	-0.0363	0.002	-15.650	0.000	-0.041
-0.032					
x4	-0.0243	0.002	-10.639	0.000	-0.029
-0.020					
x5	0.0275	0.003	9.927	0.000	0.022
0.033					
x6	0.0138	0.003	5.198	0.000	0.009
0.019					
x7	-0.0196	0.002	-8.279	0.000	-0.024
-0.015					
x8	0.1707	0.002	71.253	0.000	0.166
0.175					
x9	-0.0328	0.002	-14.029	0.000	-0.037
-0.028					
x10	-0.0252	0.003	-9.267	0.000	-0.030
-0.020					
x11	0.0896	0.002	36.281	0.000	0.085
0.094					
x12	-0.1601	0.002	-66.880	0.000	-0.165
-0.155					
x13	-0.0662	0.003	-24.345	0.000	-0.071
-0.061					
x14	-0.0462	0.002	-20.985	0.000	-0.051
-0.042					
x15	-0.0067	0.002	-3.064	0.002	-0.011
-0.002					
x16	-0.0536	0.002	-22.364	0.000	-0.058
-0.049					
x17	0.0138	0.002	5.616	0.000	0.009
0.019					
x18	-0.0108	0.003	-4.180	0.000	-0.016
-0.006					
x19	-0.0712	0.003	-28.282	0.000	-0.076
-0.066					
x20	0.0007	0.002	0.305	0.761	-0.004
0.005					
x21	0.0276	0.002	11.525	0.000	0.023
0.032					
x22	0.0557	0.003	21.536	0.000	0.051

0.061					
x23	0.0901	0.002	37.126	0.000	0.085
0.095					
x24	0.0997	0.003	33.413	0.000	0.094
0.106					
x25	0.0049	0.003	1.679	0.093	-0.001
0.011					
x26	-0.1212	0.003	-41.186	0.000	-0.127
-0.115					
x27	-0.0015	0.003	-0.523	0.601	-0.007
0.004					
x28	0.2139	0.003	71.678	0.000	0.208
0.220					
x29	0.0054	0.003	1.809	0.070	-0.000
0.011					
x30	-0.0021	0.003	-0.722	0.470	-0.008
0.004					
x31	-0.0591	0.003	-20.133	0.000	-0.065
-0.053					
x32	0.0511	0.003	17.317	0.000	0.045
0.057					
x33	0.0819	0.022	3.793	0.000	0.040
0.124					
x34	0.2518	0.046	5.474	0.000	0.162
0.342					
x35	0.1143	0.031	3.744	0.000	0.054
0.174					
x36	0.3265	0.057	5.683	0.000	0.214
0.439					
1.0/2.0	-5.7747	0.021	-272.373	0.000	-5.816
-5.733					
2.0/3.0	-0.0874	0.018	-4.879	0.000	-0.122
-0.052					
3.0/4.0	-0.2641	0.013	-20.543	0.000	-0.289
-0.239					
4.0/5.0	-0.1622	0.008	-19.599	0.000	-0.178
-0.146					
5.0/6.0	-0.0942	0.005	-17.702	0.000	-0.105
-0.084					
6.0/7.0	-0.0924	0.004	-25.372	0.000	-0.100
-0.085					
7.0/8.0	0.1040	0.002	42.815	0.000	0.099
0.109					
8.0/9.0	0.0904	0.002	40.083	0.000	0.086
0.095					
9.0/10.0	0.1427	0.003	52.318	0.000	0.137
0.148					
=====					
=====					

```

x1      0.000000e+00
x2      2.577699e-168
x3      3.321227e-55
x4      1.955738e-26
x5      3.176869e-23
x6      2.012625e-07
x7      1.246599e-16
x8      0.000000e+00
x9      1.029523e-44
x10     1.920620e-20
x11     3.235854e-288
x12     0.000000e+00
x13     6.571681e-131
x14     8.905023e-98
x15     2.182299e-03
x16     8.816428e-111
x17     1.959507e-08
x18     2.920828e-05
x19     5.724506e-176
x20     7.606672e-01
x21     9.828134e-31
x22     7.141908e-103
x23     1.054674e-301
x24     8.738269e-245
x25     9.317266e-02
x26     0.000000e+00
x27     6.010196e-01
x28     0.000000e+00
x29     7.049801e-02
x30     4.700492e-01
x31     3.813837e-90
x32     3.520680e-67
x33     1.487744e-04
x34     4.396056e-08
x35     1.808022e-04
x36     1.323051e-08
1.0/2.0  0.000000e+00
2.0/3.0  1.067125e-06
3.0/4.0  8.949131e-94
4.0/5.0  1.588290e-85
5.0/6.0  4.073080e-70
6.0/7.0  5.102751e-142
7.0/8.0  0.000000e+00
8.0/9.0  0.000000e+00
9.0/10.0 0.000000e+00
dtype: float64
      0      1
x1    -0.102683 -0.093702
x2     0.065154  0.075095

```

x3	-0.040892	-0.031789
x4	-0.028831	-0.019861
x5	0.022085	0.032951
x6	0.008585	0.018977
x7	-0.024244	-0.014962
x8	0.166049	0.175442
x9	-0.037368	-0.028207
x10	-0.030481	-0.019838
x11	0.084786	0.094470
x12	-0.164840	-0.155453
x13	-0.071488	-0.060835
x14	-0.050527	-0.041895
x15	-0.011061	-0.002431
x16	-0.058337	-0.048936
x17	0.008960	0.018567
x18	-0.015927	-0.005758
x19	-0.076156	-0.066284
x20	-0.003877	0.005304
x21	0.022886	0.032264
x22	0.050661	0.060805
x23	0.085310	0.094819
x24	0.093861	0.105559
x25	-0.000828	0.010715
x26	-0.126932	-0.115400
x27	-0.007290	0.004219
x28	0.208073	0.219772
x29	-0.000447	0.011149
x30	-0.007860	0.003626
x31	-0.064889	-0.053376
x32	0.045356	0.056933
x33	0.039603	0.124290
x34	0.161626	0.341912
x35	0.054466	0.174110
x36	0.213920	0.439149
1.0/2.0	-5.816211	-5.733103
2.0/3.0	-0.122499	-0.052284
3.0/4.0	-0.289350	-0.238945
4.0/5.0	-0.178450	-0.146003
5.0/6.0	-0.104648	-0.083784
6.0/7.0	-0.099545	-0.085268
7.0/8.0	0.099216	0.108735
8.0/9.0	0.085951	0.094788
9.0/10.0	0.137371	0.148065

It seems the accuracy is not very high. This is probably due to the difference in the number of ratings for higher values(>5) than ratings(<5)

```
import numpy as np
from sklearn.metrics import mean_squared_error, accuracy_score
```

```

# Predict probabilities on the test set
predicted_probs = model_fit.predict(X_test)

# Convert probabilities to predicted ratings by selecting the category
with the highest probability
predicted_ratings = np.argmax(predicted_probs, axis=1)

y_test = y_test.cat.codes

# Calculate evaluation metrics
mse = mean_squared_error(y_test, predicted_ratings)
accuracy = accuracy_score(y_test, predicted_ratings)

print("Mean Squared Error:", mse)
print("Accuracy:", accuracy)

Mean Squared Error: 2.736833967282061
Accuracy: 0.260828153508365

import matplotlib.pyplot as plt
import seaborn as sns


plt.figure(figsize=(10, 6))
sns.scatterplot(x=y_test, y=predicted_ratings, alpha=0.5)

plt.plot([1, 10], [1, 10], color='red', linestyle='--')

plt.title('Actual vs. Predicted Ratings')
plt.xlabel('Actual Ratings')
plt.ylabel('Predicted Ratings')
plt.xlim(1, 10)
plt.ylim(1, 10)
plt.xticks(range(1, 11))
plt.yticks(range(1, 11))

# Show plot
plt.grid()
plt.show()

```



Undersampling the data based on rating to reduce the difference in number of ratings

```
rating_counts = df['rating'].value_counts()

target_sample_size = rating_counts.nlargest(5).min()

rating_under_sampled_df = df.groupby('rating').apply(lambda x:
x.sample(n=min(len(x), target_sample_size),
random_state=42)).reset_index(drop=True)

target_counts = rating_under_sampled_df['rating'].value_counts()

print("Under-sampled DataFrame shape:", rating_under_sampled_df.shape)
print("Counts of each class in the target variable:")
print(target_counts)
```

```
Under-sampled DataFrame shape: (521841, 51)
Counts of each class in the target variable:
rating
6.0    88831
7.0    88831
8.0    88831
9.0    88831
```

```
10.0    88831
5.0     44137
4.0     18756
3.0       7887
2.0       4136
1.0       2770
```

```
Name: count, dtype: int64
```

```
/var/folders/dz/fg9tl53x4y16ytmghwdt0kr0000gn/T/
```

```
ipykernel_85581/1598312855.py:7: DeprecationWarning:
```

```
DataFrameGroupBy.apply operated on the grouping columns. This behavior
is deprecated, and in a future version of pandas the grouping columns
will be excluded from the operation. Either pass
`include_groups=False` to exclude the groupings or explicitly select
the grouping columns after groupby to silence this warning.
```

```
rating_under_sampled_df = df.groupby('rating').apply(lambda x:
x.sample(n=min(len(x), target_sample_size),
random_state=42)).reset_index(drop=True)
```

Second iteration with undersampled rating data

```
# Standardize the features
```

```
features = rating_under_sampled_df.columns.difference(['user_id',
'anime_id', 'rating', 'Birthday_Date', 'Joined_Date', 'Age_Join',
'Episodes Watched',
                                                    'Start Date', 'End Date', 'Name',
'Rank', 'Studios', 'Source', 'Episodes_Norm', 'Unnamed: 0'])
```

```
# # Split the data
```

```
X = rating_under_sampled_df[features]
```

```
y = rating_under_sampled_df[target]
```

```
# Ordinal Logistic Regression expects integer targets as levels
```

```
y = y.astype('category')
```

```
# # Split into training and test sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

```
# Fit an Ordinal Logistic Regression model
```

```
model_under_rating = OrderedModel(y_train, X_train, distr='logit')
```

```
model_under_rating_fit = model_under_rating.fit(method='bfgs')
```

```
print(model_under_rating_fit.summary())
```

```
# Coefficient significance
```

```
print(model_under_rating_fit.pvalues)
```

```
print(model_under_rating_fit.conf_int())
```

```
/Users/ramachandrak/Repos/MS/sem1/.venv/lib/python3.12/site-  
packages/statsmodels/miscmodels/ordinal_model.py:205: Warning: the  
endog has ordered == False, risk of capturing a wrong order for the  
categories. ordered == True preferred.
```

```
warnings.warn("the endog has ordered == False, "
```

```
Optimization terminated successfully.
```

```
Current function value: 1.921322
```

```
Iterations: 205
```

```
Function evaluations: 210
```

```
Gradient evaluations: 210
```

```
OrderedModel Results
```

```
=====
```

```
Dep. Variable:          rating    Log-Likelihood:      -  
8.0210e+05
```

```
Model:          OrderedModel    AIC:  
1.604e+06
```

```
Method:          Maximum Likelihood    BIC:  
1.605e+06
```

```
Date:          Tue, 05 Nov 2024
```

```
Time:          02:19:00
```

```
No. Observations:          417472
```

```
Df Residuals:          417427
```

```
Df Model:          36
```

```
=====
```

```
=====
```

```
coef    std err          z    P>|z|
```

```
[0.025    0.975]
```

```
-----
```

```
Age          -0.0247    0.001   -36.055    0.000
```

```
-0.026    -0.023
```

```
Episodes          0.0015    6.44e-05    23.744    0.000
```

```
0.001    0.002
```

```
Gender_Male     -0.0734    0.006   -12.010    0.000
```

```
-0.085    -0.061
```

```
Gender_Non-Binary -0.3315    0.033    -9.961    0.000
```

```
-0.397    -0.266
```

```
Genre_Action          0.0646    0.007     9.348    0.000
```

```
0.051    0.078
```

```
Genre_Adventure          0.0407    0.008     4.906    0.000
```

```
0.024    0.057
```

```
Genre_Avant Garde   -0.2243    0.032    -7.093    0.000
```



-0.286	-0.162				
Genre_Award Winning	0.856	0.8839	0.014	62.960	0.000
	0.911				
Genre_Boys Love	-0.476	-0.4206	0.028	-14.848	0.000
	-0.365				
Genre_Comedy	-0.050	-0.0372	0.007	-5.471	0.000
	-0.024				
Genre_Drama	0.196	0.2091	0.007	30.387	0.000
	0.223				
Genre_Ecchi	-0.546	-0.5282	0.009	-59.781	0.000
	-0.511				
Genre_Fantasy	-0.182	-0.1668	0.008	-21.968	0.000
	-0.152				
Genre_Girls Love	-0.507	-0.4568	0.025	-17.968	0.000
	-0.407				
Genre_Gourmet	-0.132	-0.0705	0.031	-2.259	0.024
	-0.009				
Genre_Horror	-0.304	-0.2788	0.013	-21.321	0.000
	-0.253				
Genre_Mystery	0.026	0.0442	0.009	4.728	0.000
	0.062				
Genre_Romance	-0.033	-0.0192	0.007	-2.782	0.005
	-0.006				
Genre_Sci-Fi	-0.212	-0.1965	0.008	-25.039	0.000
	-0.181				
Genre_Slice of Life	-0.002	0.0237	0.013	1.840	0.066
	0.049				
Genre_Sports	0.141	0.1740	0.017	10.308	0.000
	0.207				
Genre_Supernatural	0.134	0.1495	0.008	18.768	0.000
	0.165				
Genre_Suspense	0.462	0.4901	0.014	34.748	0.000
	0.518				
Location_Brazil	0.359	0.3828	0.012	31.132	0.000
	0.407				
Location_Canada	0.041	0.0645	0.012	5.293	0.000
	0.088				
Location_France	-0.437	-0.4132	0.012	-34.098	0.000
	-0.389				
Location_Germany	-0.022	0.0018	0.012	0.148	0.882
	0.026				
Location_Philippines	0.819	0.8435	0.012	68.816	0.000
	0.867				
Location_Poland	0.015	0.0393	0.012	3.194	0.001
	0.063				
Location_Russia	-0.020	0.0039	0.012	0.316	0.752
	0.028				
Location_Sweden	-0.206	-0.1822	0.012	-15.093	0.000
	-0.159				

Location_United States	0.2110	0.012	17.251	0.000
0.187 0.235				
Type_ONA	0.3814	0.200	1.907	0.057
-0.011 0.773				
Type_OVA	0.7052	0.199	3.545	0.000
0.315 1.095				
Type_Special	0.3728	0.199	1.871	0.061
-0.018 0.763				
Type_TV	0.7456	0.199	3.753	0.000
0.356 1.135				
1.0/2.0	-5.3759	0.201	-26.723	0.000
-5.770 -4.982				
2.0/3.0	-0.0690	0.018	-3.860	0.000
-0.104 -0.034				
3.0/4.0	-0.2473	0.013	-19.321	0.000
-0.272 -0.222				
4.0/5.0	-0.1396	0.008	-17.000	0.000
-0.156 -0.123				
5.0/6.0	-0.0503	0.005	-9.572	0.000
-0.061 -0.040				
6.0/7.0	0.0244	0.004	6.893	0.000
0.017 0.031				
7.0/8.0	-0.2750	0.003	-78.730	0.000
-0.282 -0.268				
8.0/9.0	-0.2826	0.003	-80.926	0.000
-0.289 -0.276				
9.0/10.0	-0.0231	0.004	-6.544	0.000
-0.030 -0.016				
=====				
=====				
Age	1.169237e-284			
Episodes	1.257567e-124			
Gender_Male	3.141376e-33			
Gender_Non-Binary	2.249364e-23			
Genre_Action	8.926337e-21			
Genre_Adventure	9.297422e-07			
Genre_Avant Garde	1.315083e-12			
Genre_Award Winning	0.000000e+00			
Genre_Boys Love	7.129536e-50			
Genre_Comedy	4.486878e-08			
Genre_Drama	8.231411e-203			
Genre_Ecchi	0.000000e+00			
Genre_Fantasy	5.862816e-107			
Genre_Girls Love	3.448285e-72			
Genre_Gourmet	2.388877e-02			
Genre_Horror	7.316307e-101			
Genre_Mystery	2.267776e-06			
Genre_Romance	5.401458e-03			
Genre_Sci-Fi	2.277079e-138			

Genre_Slice of Life	6.571469e-02
Genre_Sports	6.507695e-25
Genre_Supernatural	1.381168e-78
Genre_Suspense	1.469830e-264
Location_Brazil	8.905444e-213
Location_Canada	1.204592e-07
Location_France	7.868142e-255
Location_Germany	8.824292e-01
Location_Philippines	0.000000e+00
Location_Poland	1.402949e-03
Location_Russia	7.523721e-01
Location_Sweden	1.794018e-51
Location_United States	1.107131e-66
Type_ONA	5.654060e-02
Type_OVA	3.923281e-04
Type_Special	6.136839e-02
Type_TV	1.747418e-04
1.0/2.0	2.527731e-157
2.0/3.0	1.135987e-04
3.0/4.0	3.603350e-83
4.0/5.0	8.244007e-65
5.0/6.0	1.053507e-21
6.0/7.0	5.466686e-12
7.0/8.0	0.000000e+00
8.0/9.0	0.000000e+00
9.0/10.0	5.998021e-11

dtype: float64

	0	1
Age	-0.026047	-0.023361
Episodes	0.001404	0.001657
Gender_Male	-0.085393	-0.061432
Gender_Non-Binary	-0.396676	-0.266242
Genre_Action	0.051034	0.078111
Genre_Adventure	0.024451	0.056985
Genre_Avant Garde	-0.286247	-0.162298
Genre_Award Winning	0.856408	0.911442
Genre_Boys Love	-0.476156	-0.365110
Genre_Comedy	-0.050486	-0.023852
Genre_Drama	0.195632	0.222609
Genre_Ecchi	-0.545526	-0.510891
Genre_Fantasy	-0.181644	-0.151887
Genre_Girls Love	-0.506618	-0.406965
Genre_Gourmet	-0.131694	-0.009332
Genre_Horror	-0.304466	-0.253201
Genre_Mystery	0.025864	0.062492
Genre_Romance	-0.032701	-0.005669
Genre_Sci-Fi	-0.211833	-0.181078
Genre_Slice of Life	-0.001538	0.048861
Genre_Sports	0.140917	0.207090

Genre_Supernatural	0.133878	0.165101
Genre_Suspense	0.462463	0.517751
Location_Brazil	0.358706	0.406906
Location_Canada	0.040638	0.088434
Location_France	-0.436949	-0.389447
Location_Germany	-0.022208	0.025833
Location_Philippines	0.819436	0.867482
Location_Poland	0.015194	0.063458
Location_Russia	-0.020302	0.028092
Location_Sweden	-0.205896	-0.158568
Location_United States	0.187005	0.234945
Type_ONA	-0.010624	0.773351
Type_OVA	0.315323	1.095049
Type_Special	-0.017761	0.763361
Type_TV	0.356205	1.134945
1.0/2.0	-5.770132	-4.981570
2.0/3.0	-0.104109	-0.033983
3.0/4.0	-0.272350	-0.222182
4.0/5.0	-0.155652	-0.123471
5.0/6.0	-0.060540	-0.039960
6.0/7.0	0.017459	0.031332
7.0/8.0	-0.281800	-0.268110
8.0/9.0	-0.289480	-0.275790
9.0/10.0	-0.029995	-0.016169

```
# Predict probabilities on the test set
```

```
predicted_probs = model_under_rating_fit.predict(X_test)
```

```
# Convert probabilities to predicted ratings by selecting the category with the highest probability
```

```
predicted_ratings = np.argmax(predicted_probs, axis=1)
```

```
y_test = y_test.astype('category').cat.codes + 1
```

```
# Calculate evaluation metrics
```

```
mse = mean_squared_error(y_test, predicted_ratings)
```

```
accuracy = accuracy_score(y_test, predicted_ratings)
```

```
print("Mean Squared Error:", mse)
```

```
print("Accuracy:", accuracy)
```

```
Mean Squared Error: 5.318926117908575
```

```
Accuracy: 0.163640544606157
```

```
plt.figure(figsize=(10, 6))
```


```
sns.scatterplot(x=y_test, y=predicted_ratings, alpha=0.5)
```

```
plt.plot([1, 10], [1, 10], color='red', linestyle='--')
```

```
plt.title('Actual vs. Predicted Ratings')
```

```
plt.xlabel('Actual Ratings')
plt.ylabel('Predicted Ratings')
plt.xlim(1, 10)
plt.ylim(1, 10)
plt.xticks(range(1, 11))
plt.yticks(range(1, 11))

# Show plot
plt.grid()
plt.show()
```



```
import pickle

# Save the model to a file
with open('./models/ordinal_logistic_model.pkl', 'wb') as f:
    pickle.dump(model_fit, f)

# Standardize the features
features = rating_under_sampled_df.columns.difference(['user_id',
'anime_id', 'rating', 'Birthday_Date', 'Joined_Date', 'Age_Join',
'Episodes Watched',
'Start Date', 'End Date', 'Name',
'Rank', 'Studios', 'Source', 'Episodes_Norm', 'Unnamed: 0'])
from sklearn.preprocessing import LabelEncoder, StandardScaler
```

```
scaler = StandardScaler()
df_scaled = scaler.fit_transform(rating_under_sampled_df[features])
```

```
# # Split the data
```

```
X = df_scaled
```

```
y = rating_under_sampled_df[target]
```

```
# Ordinal Logistic Regression expects integer targets as levels
```

```
y = y.astype('category')
```

```
# # Split into training and test sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

```
# Fit an Ordinal Logistic Regression model
```

```
model_under_rating_standard = OrderedModel(y_train, X_train,
distr='logit')
```

```
model_under_rating_standard_fit =
```

```
model_under_rating_standard.fit(method='bfgs')
```

```
print(model_under_rating_standard_fit.summary())
```

```
# Coefficient significance
```

```
print(model_under_rating_standard_fit.pvalues)
```

```
print(model_under_rating_standard_fit.conf_int())
```

```
/Users/ramachandrank/Repos/MS/sem1/.venv/lib/python3.12/site-
packages/statsmodels/miscmodels/ordinal_model.py:205: Warning: the
endog has ordered == False, risk of capturing a wrong order for the
categories. ordered == True preferred.
```

```
warnings.warn("the endog has ordered == False, "
```

```
Optimization terminated successfully.
```

```
Current function value: 1.921322
```

```
Iterations: 70
```

```
Function evaluations: 71
```

```
Gradient evaluations: 71
```

```
OrderedModel Results
```

```
=====
=====
```

```
Dep. Variable:                rating    Log-Likelihood:    -
8.0210e+05
```

```
Model:                        OrderedModel    AIC:
```

```
1.604e+06
```

```
Method:                      Maximum Likelihood    BIC:
```

```
1.605e+06
```

```
Date:                        Tue, 05 Nov 2024
```

```
Time:                        02:29:34
```

No. Observations: 417472

Df Residuals: 417427

Df Model: 36

=====

=====

coef std err z P>|z| [0.025

0.975]

-----

x1 -0.1021 0.003 -36.055 0.000 -0.108

-0.097

x2 0.0753 0.003 23.745 0.000 0.069

0.082

x3 -0.0345 0.003 -12.010 0.000 -0.040

-0.029

x4 -0.0279 0.003 -9.960 0.000 -0.033

-0.022

x5 0.0320 0.003 9.345 0.000 0.025

0.039

x6 0.0161 0.003 4.891 0.000 0.010

0.023

x7 -0.0209 0.003 -7.099 0.000 -0.027

-0.015

x8 0.1884 0.003 62.968 0.000 0.182

0.194

x9 -0.0424 0.003 -14.848 0.000 -0.048

-0.037

x10 -0.0184 0.003 -5.476 0.000 -0.025

-0.012

x11 0.0930 0.003 30.385 0.000 0.087

0.099

x12 -0.1769 0.003 -59.767 0.000 -0.183

-0.171

x13 -0.0738 0.003 -21.956 0.000 -0.080

-0.067

x14 -0.0488 0.003 -17.949 0.000 -0.054

-0.043

x15 -0.0061 0.003 -2.252 0.024 -0.011

-0.001

x16 -0.0632 0.003 -21.315 0.000 -0.069

-0.057

x17 0.0143 0.003 4.736 0.000 0.008

0.020

x18 -0.0090 0.003 -2.788 0.005 -0.015

-0.003

x19 -0.0778 0.003 -25.039 0.000 -0.084

-0.072					
x20	0.0053	0.003	1.839	0.066	-0.000
0.011					
x21	0.0304	0.003	10.308	0.000	0.025
0.036					
x22	0.0601	0.003	18.767	0.000	0.054
0.066					
x23	0.1048	0.003	34.745	0.000	0.099
0.111					
x24	0.1159	0.004	31.156	0.000	0.109
0.123					
x25	0.0194	0.004	5.306	0.000	0.012
0.027					
x26	-0.1250	0.004	-34.100	0.000	-0.132
-0.118					
x27	0.0006	0.004	0.154	0.877	-0.007
0.008					
x28	0.2550	0.004	68.830	0.000	0.248
0.262					
x29	0.0118	0.004	3.211	0.001	0.005
0.019					
x30	0.0012	0.004	0.323	0.747	-0.006
0.008					
x31	-0.0551	0.004	-15.070	0.000	-0.062
-0.048					
x32	0.0632	0.004	17.261	0.000	0.056
0.070					
x33	0.0441	0.023	1.912	0.056	-0.001
0.089					
x34	0.1724	0.049	3.553	0.000	0.077
0.267					
x35	0.0612	0.033	1.876	0.061	-0.003
0.125					
x36	0.2288	0.061	3.760	0.000	0.110
0.348					
1.0/2.0	-5.3784	0.021	-250.668	0.000	-5.420
-5.336					
2.0/3.0	-0.0688	0.018	-3.844	0.000	-0.104
-0.034					
3.0/4.0	-0.2474	0.013	-19.327	0.000	-0.272
-0.222					
4.0/5.0	-0.1395	0.008	-16.998	0.000	-0.156
-0.123					
5.0/6.0	-0.0502	0.005	-9.570	0.000	-0.061
-0.040					
6.0/7.0	0.0244	0.004	6.890	0.000	0.017
0.031					
7.0/8.0	-0.2749	0.003	-78.727	0.000	-0.282
-0.268					



8.0/9.0	-0.2827	0.003	-80.933	0.000	-0.290
-0.276					
9.0/10.0	-0.0231	0.004	-6.539	0.000	-0.030
-0.016					

=====

=====

x1	1.137272e-284
x2	1.240609e-124
x3	3.134745e-33
x4	2.278068e-23
x5	9.215770e-21
x6	1.004153e-06
x7	1.259340e-12
x8	0.000000e+00
x9	7.127258e-50
x10	4.350891e-08
x11	8.616374e-203
x12	0.000000e+00
x13	7.632627e-107
x14	4.921900e-72
x15	2.433157e-02
x16	8.164416e-101
x17	2.175851e-06
x18	5.308597e-03
x19	2.286305e-138
x20	6.585217e-02
x21	6.452772e-25
x22	1.411825e-78
x23	1.635496e-264
x24	4.265630e-213
x25	1.121340e-07
x26	7.419827e-255
x27	8.772316e-01
x28	0.000000e+00
x29	1.323391e-03
x30	7.470042e-01
x31	2.564371e-51
x32	9.237565e-67
x33	5.588277e-02
x34	3.814004e-04
x35	6.067350e-02
x36	1.696355e-04
1.0/2.0	0.000000e+00
2.0/3.0	1.208487e-04
3.0/4.0	3.175558e-83
4.0/5.0	8.549532e-65
5.0/6.0	1.064837e-21
6.0/7.0	5.593449e-12
7.0/8.0	0.000000e+00

8.0/9.0        0.000000e+00

9.0/10.0      6.177096e-11

dtype: float64

	0	1
x1	-0.107699	-0.096594
x2	0.069124	0.081562
x3	-0.040172	-0.028900
x4	-0.033382	-0.022404
x5	0.025281	0.038700
x6	0.009628	0.022505
x7	-0.026611	-0.015096
x8	0.182493	0.194219
x9	-0.048000	-0.036806
x10	-0.024964	-0.011804
x11	0.086993	0.098990
x12	-0.182710	-0.171107
x13	-0.080335	-0.067168
x14	-0.054111	-0.043457
x15	-0.011492	-0.000796
x16	-0.068995	-0.057376
x17	0.008400	0.020260
x18	-0.015249	-0.002659
x19	-0.083943	-0.071755
x20	-0.000349	0.011009
x21	0.024636	0.036203
x22	0.053794	0.066341
x23	0.098892	0.110716
x24	0.108634	0.123220
x25	0.012221	0.026538
x26	-0.132231	-0.117857
x27	-0.006577	0.007703
x28	0.247776	0.262301
x29	0.004607	0.019044
x30	-0.005973	0.008326
x31	-0.062240	-0.047913
x32	0.056018	0.070370
x33	-0.001109	0.089405
x34	0.077282	0.267490
x35	-0.002742	0.125080
x36	0.109527	0.347990
1.0/2.0	-5.420454	-5.336347
2.0/3.0	-0.103820	-0.033706
3.0/4.0	-0.272443	-0.222274
4.0/5.0	-0.155633	-0.123452
5.0/6.0	-0.060534	-0.039954
6.0/7.0	0.017447	0.031321
7.0/8.0	-0.281789	-0.268099
8.0/9.0	-0.289509	-0.275818
9.0/10.0	-0.029980	-0.016153

```

# Predict probabilities on the test set
predicted_probs = model_under_rating_standard_fit.predict(X_test)

# Convert probabilities to predicted ratings by selecting the category
with the highest probability
predicted_ratings = np.argmax(predicted_probs, axis=1) + 1

y_test = y_test.astype('category').cat.codes + 1

# Calculate evaluation metrics
mse = mean_squared_error(y_test, predicted_ratings)
accuracy = accuracy_score(y_test, predicted_ratings)

print("Mean Squared Error:", mse)
print("Accuracy:", accuracy)

Mean Squared Error: 5.318926117908575
Accuracy: 0.163640544606157


plt.figure(figsize=(10, 6))
sns.scatterplot(x=y_test, y=predicted_ratings, alpha=0.5)

plt.plot([1, 10], [1, 10], color='red', linestyle='--')

plt.title('Actual vs. Predicted Ratings')
plt.xlabel('Actual Ratings')
plt.ylabel('Predicted Ratings')
plt.xlim(1, 10)
plt.ylim(1, 10)
plt.xticks(range(1, 11))
plt.yticks(range(1, 11))

# Show plot
plt.grid()
plt.show()

```



Ordered Model was not very accurate with our data. This is mostly due to the sparseness of our data after 1 hot encoding multiple class columns like location and genres.