

Rutuja Badve First we will load our cleaned datasets

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
anime_df = pd.read_csv("anime_dataset.xls")
user_scores = pd.read_csv("user_scores_d.xls")
users_details = pd.read_csv("users_details_dataset.xls")
```

```
anime_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19976 entries, 0 to 19975
Data columns (total 29 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            19976 non-null  int64
1   anime_id              19976 non-null  int64
2   Name                  19976 non-null  object
3   English name          19976 non-null  object
4   Other name            19976 non-null  object
5   Score                 19976 non-null  object
6   Genres                 19976 non-null  object
7   Synopsis              19976 non-null  object
8   Type                  19976 non-null  object
9   Episodes              19682 non-null  float64
10  Aired                  19976 non-null  object
11  Start Date            16957 non-null  object
12  End Date              8858 non-null  object
13  Premiered             19976 non-null  object
14  Status                 19976 non-null  object
15  Producers              19976 non-null  object
16  Licensors              19976 non-null  object
17  Studios                19976 non-null  object
18  Source                 19976 non-null  object
19  Duration               19976 non-null  object
20  Rating                 19976 non-null  object
21  Rank                   19976 non-null  object
22  Popularity             19976 non-null  int64
23  Favorites              19976 non-null  int64
24  Scored By              19976 non-null  object
25  Members                19976 non-null  int64
26  Image URL              19976 non-null  object
27  Ongoing                19976 non-null  int64
28  Episodes_Norm          19682 non-null  float64
dtypes: float64(2), int64(6), object(21)
memory usage: 4.4+ MB
```

```
users_details.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 41486 entries, 0 to 41485
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            41486 non-null  int64
1   Mal ID                 41486 non-null  int64
2   Username                41486 non-null  object
3   Gender                  40316 non-null  object
4   Birthday                41486 non-null  object
5   Location                41486 non-null  object
6   Joined                  41486 non-null  object
7   Days Watched            41485 non-null  float64
8   Mean Score              41485 non-null  float64
9   Watching                41485 non-null  float64
10  Completed               41485 non-null  float64
11  On Hold                  41485 non-null  float64
12  Dropped                  41485 non-null  float64
13  Plan to Watch            41485 non-null  float64
14  Total Entries            41485 non-null  float64
15  Rewatched                41485 non-null  float64
16  Episodes Watched         41485 non-null  float64
17  Birthday_Date            41486 non-null  object
18  Joined_Date              41486 non-null  object
19  Age_Join                 41486 non-null  float64
dtypes: float64(11), int64(2), object(7)
memory usage: 6.3+ MB
```

```
user_scores.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 523975 entries, 0 to 523974
Data columns (total 6 columns):
```

```

#    Column      Dtype
---  -
0    Unnamed: 0    int64
1    user_id       int64
2    Username      object
3    anime_id      int64
4    Anime Title   object
5    rating        int64
dtypes: int64(4), object(2)
memory usage: 239.9+ MB

```

1. Do different age groups prefer anime with varying durations?

Ans: We are using K means clustering to find the relationship between age groups and duration of episodes of animes.

```

joined_df = users_details.merge(user_scores, on='Username')
joined_df = joined_df.merge(anime_df[['anime_id', 'Episodes']], on='anime_id')
joined_df = joined_df.dropna(subset=['Age_Join', 'Episodes'])

```

```

from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans

```

```
Feature = joined_df[['Age_Join', 'Episodes']]
```

```

scaler = StandardScaler()
Feature_scaled = scaler.fit_transform(Feature)

```

```

kmeans = KMeans(n_clusters=6, random_state=8)
joined_df['Cluster'] = kmeans.fit_predict(Feature_scaled)

```

After trying for different k values, have selected 6 as it is giving classification which includes variance in our dataset.

I have chosen K means for following reasons:

- As we want to cluster unknown data, as we are not familiar with relationship of age group and number of episodes we use K means because it is a simple clustering algorithm.
- As it is unsupervised algorithm it will be able to find the hidden relationship among parameters
- K means works well with large datasets and our dataset is very large


For training the Model:

- Have selected the important features like age and number of episodes which are necessary
- Have normalized the data to fit well for our prediction
- The number of k has been chosen by trial and error method

```

plt.figure(figsize=(10, 7))
plt.scatter(joined_df['Age_Join'], joined_df['Episodes'], c=joined_df['Cluster'], cmap='viridis')
plt.xlabel('Age')
plt.ylabel('Number of Episodes')
plt.title('K-means Clustering of Users by Age and Episodes Watched')
plt.colorbar(label='Cluster')
plt.show()

```



The K means have given us well defined clusters and we can make out some predictions through the above graph

We can see that the older people do not prefer animes which have lots of episodes.

There is a age group between 10 to 35 which can watch many number of episodes but the same group also has sub groups within it that prefer watching different length of anime

Now we will use DBSCAN Algorithm to analyze the first question

Reference taken

from: <https://builtin.com/articles/dbscan#:~:text=DBSCAN%20works%20by%20partitioning%20the,are%20considered%20outliers%20or%20noise>

We are using DBSCAN as it is better than K-means for clustering purpose, K-means will have only sperical clusters but DBSCAN can have various shaped clusters. Here as our data and relationship is unknown DBSCAN can give us more accurate results.

It is also useful for identifying outliers, the gray points that we observe in output plot are the outliers and we can get rid of them to make our model more accurate.

DBSCAN gave us 7 different clusters indicating that there exists 7 different classes in our dataset, we can use this information in our K- means algorithm.

```
#https://builtin.com/articles/dbscan#:~:text=DBSCAN%20works%20by%20partitioning%20the,are%20considered%20outliers%20or%20noise.
```

```
from sklearn.cluster import DBSCAN
sampled_df = Feature.sample(frac=0.4, random_state=42)
```

```
sampled_df =sampled_df[['Age_Join', 'Episodes']]
```

```
scaler = StandardScaler()
scaled_data = scaler.fit_transform(sampled_df)
```

```
dbscan = DBSCAN(eps=0.5, min_samples=5)
labels = dbscan.fit_predict(scaled_data)
```

```
n_clusters = len(set(labels)) - (1 if -1 in labels else 0)
print(f"Number of clusters found by DBSCAN: {n_clusters}")
```

```
unique_labels = set(labels)
colors = [plt.cm.Spectral(each) for each in np.linspace(0, 1, len(unique_labels))]
```

```
plt.figure(figsize=(10, 6))
for k, col in zip(unique_labels, colors):
    if k == -1:
        col = [0, 0, 0, 1]
```

```

class_member_mask = (labels == k)
xy = scaled_data[class_member_mask]
plt.scatter(xy[:, 0], xy[:, 1], s=50, c=[col], marker='o', alpha=0.5)

plt.title('DBSCAN Clustering on Age and Episodes')
plt.xlabel('Age_scaled')
plt.ylabel('Episodes_scaled')
plt.show()

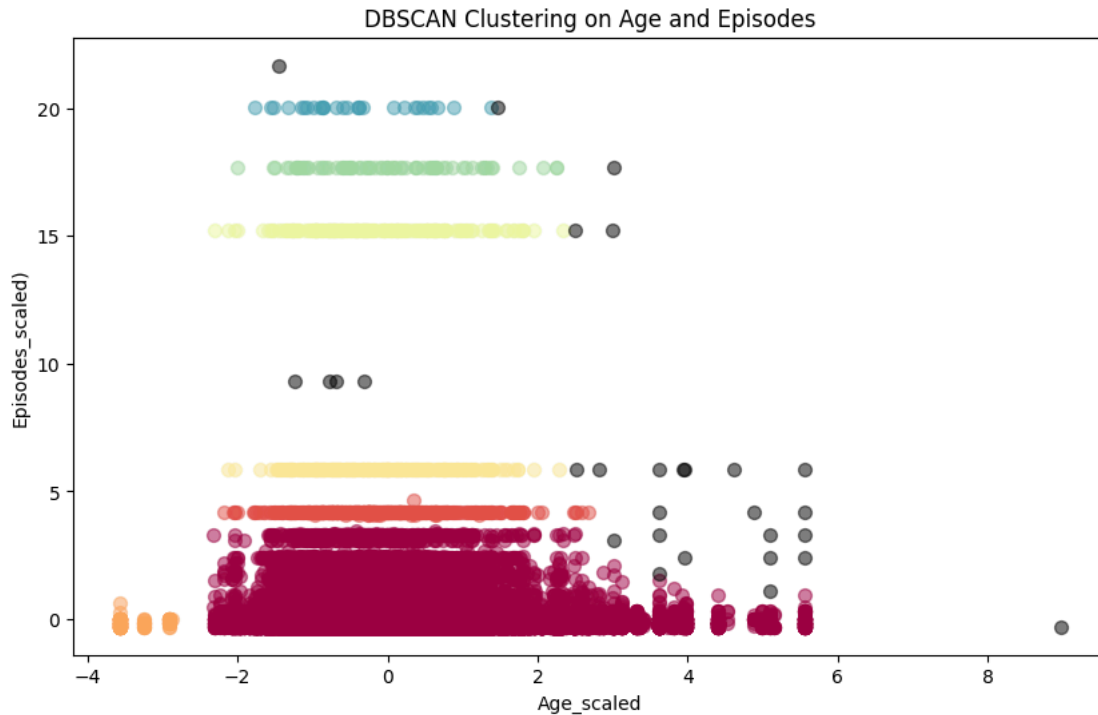
```

```

print("Indices of core samples: ", dbscan.core_sample_indices_)
print("Core samples_centroids: \n", dbscan.components_)
print("Labels assigned to each point:", dbscan.labels_)
print("Number of features seen during fit: ", dbscan.n_features_in_)

```

→ Number of clusters found by DBSCAN: 7



```

Indices of core samples: [ 0 1 2 ... 117539 117540 117541]
Core samples_centroids:
[[ 0.03212692 -0.2887712 ]
 [ 0.12026636  0.01860741]
 [-0.67138602 -0.2887712 ]
 ...
 [ 2.29170151 -0.14122947]
 [ 0.18917537 -0.2887712 ]
 [ 0.51288747 -0.27647606]]
Labels assigned to each point: [0 0 0 ... 0 0 0]
Number of features seen during fit: 2

```

2.Does the number of episodes affect the score of anime?

Ans: For this we first plot a graph to see if the relationship between score and number of episodes is linear or not.

```

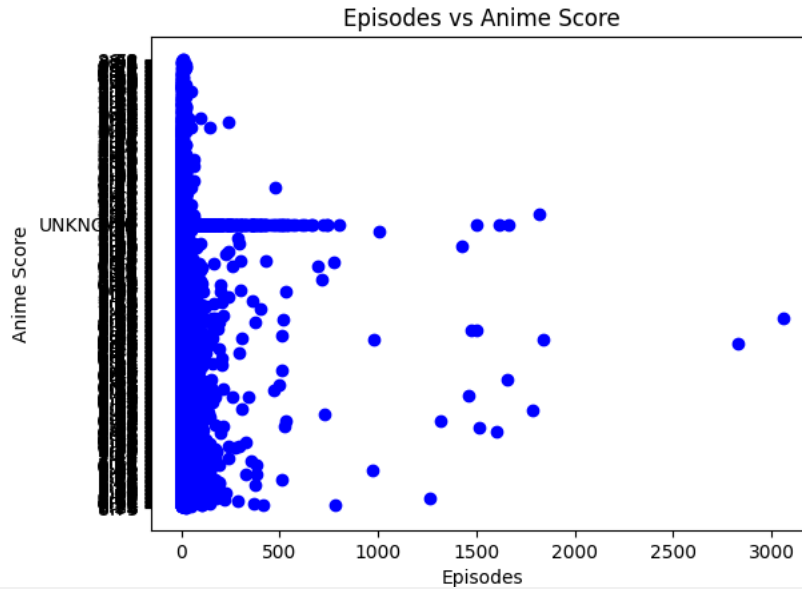
#We plot the relationship to find out if it is linear or not
import matplotlib.pyplot as plt

plt.scatter(anime_df['Episodes'], anime_df['Score'], color='blue')

plt.title('Episodes vs Anime Score')
plt.xlabel('Episodes')
plt.ylabel('Anime Score')

```

```
Text(0, 0.5, 'Anime Score')
```



As the relationship is not linear we use non linear regression algorithms

We are using Support Vector Regression algorithm: Reference taken from: <https://www.geeksforgeeks.org/support-vector-regression-svr-using-linear-and-non-linear-kernels-in-scikit-learn/>

SVR is robust to outliers and as we use it in large dataset we are removing the outliers.

```
#https://www.geeksforgeeks.org/support-vector-regression-svr-using-linear-and-non-linear-kernels-in-scikit-learn/
from sklearn.svm import SVR
```

```
anime_df['Episodes'] = pd.to_numeric(anime_df['Episodes'], errors='coerce')
anime_df['Score'] = pd.to_numeric(anime_df['Score'], errors='coerce')
```

```
anime_df = anime_df.dropna(subset=['Episodes', 'Score'])
```


```
X = anime_df[['Episodes']].values
y = anime_df['Score'].values
```

```
svr = SVR(kernel='rbf')
```

```
svr.fit(X, y)
```

```
y_pred = svr.predict(X)
```

```
plt.scatter(X, y, color='darkorange', label='Actual Data')
plt.plot(X, y_pred, color='cornflowerblue', label='SVR Prediction')
plt.xlabel('Episodes')
plt.ylabel('Anime Score')
plt.title('Support Vector Regression: Episodes vs. Anime Score')
plt.legend()
plt.show()
```



```
from sklearn.metrics import mean_absolute_error, mean_squared_error
mae = mean_absolute_error(y, y_pred)
mse = mean_squared_error(y, y_pred)
```

```
print(f"Mean Absolute Error (MAE): {mae}")
print(f"Mean Squared Error (MSE): {mse}")
```

```
Mean Absolute Error (MAE): 0.7019815620007238
Mean Squared Error (MSE): 0.7727067390175613
```

SVR gave us moderate results to improve them we need to hypertune our parameters

We are using XGBoost Algorithm: Reference taken from: <https://www.datacamp.com/tutorial/xgboost-in-python>

XGBoost excels in capturing complex relationships

```
!pip install xgboost
```

```
Collecting xgboost
  Downloading xgboost-2.1.2-py3-none-manylinux_2_28_x86_64.whl.metadata (2.1 kB)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from xgboost) (1.26.4)
Collecting nvidia-nccl-cu12 (from xgboost)
  Downloading nvidia_nccl_cu12-2.23.4-py3-none-manylinux2014_x86_64.whl.metadata (1.8 kB)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from xgboost) (1.13.1)
Downloading xgboost-2.1.2-py3-none-manylinux_2_28_x86_64.whl (153.9 MB)
 153.9/153.9 MB 7.0 MB/s eta 0:00:00
Downloading nvidia_nccl_cu12-2.23.4-py3-none-manylinux2014_x86_64.whl (199.0 MB)
 199.0/199.0 MB 5.3 MB/s eta 0:00:00
Installing collected packages: nvidia-nccl-cu12, xgboost
Successfully installed nvidia-nccl-cu12-2.23.4 xgboost-2.1.2
```

```
import xgboost as xgb
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import LabelEncoder

anime_df = anime_df.dropna(subset=['Episodes', 'Score'])

anime_df_sampled = anime_df.sample(frac=0.3, random_state=42)

cat_cols = anime_df_sampled.select_dtypes(include=['object']).columns.tolist()

label_encoder = LabelEncoder()

for col in cat_cols:
    anime_df_sampled[col] = label_encoder.fit_transform(anime_df_sampled[col])

X = anime_df_sampled[['Episodes']]
y = anime_df_sampled['Score']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)

dtrain_xg = xgb.DMatrix(X_train, y_train)
```

```
dtest_xg = xgb.DMatrix(X_test, y_test)
```

```
params = {
    "objective": "reg:squarederror",
    "eval_metric": "rmse"
}
```

```
n_rounds = 180
```

```
model = xgb.train(
    params=params,
    dtrain=dtrain_xg,
    num_boost_round=n_rounds
)
```

```
preds = model.predict(dtest_xg)
```

```
rmse = mean_squared_error(y_test, preds, squared=False)
print(f"RMSE of the XGBoost model: {rmse:.3f}")
```

```
→ RMSE of the XGBoost model: 126.593
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_regression.py:492: FutureWarning: 'squared' is deprecated in version 1.4 ar
warnings.warn(
```

```
print(y.describe())
```

```
→ count    3936.000000
   mean      306.004573
   std      125.746713
   min         0.000000
   25%      209.000000
   50%      300.000000
   75%      461.000000
   max      461.000000
   Name: Score, dtype: float64
```

```
import matplotlib.pyplot as plt
```

```
plt.figure(figsize=(8, 7))
plt.scatter(y_test, preds, color='blue', alpha=0.6)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'k--', lw=2)
plt.xlabel('Actual Score')
plt.ylabel('Predicted_Score')
plt.title('Predicted vs Actual Score')
plt.show()
```